

UNIT-3

process Synchronization

Our original solution allowed at most $\text{BUFSIZE} - 1$ items in the buffer at the same time.

Suppose we want to modify the algorithm to remedy this deficiency. One possibility is to add an integer variable counter, initialized to 0. Counter is incremented every time we add a new item to the buffer.

→ It is decremented every time we remove one item from the buffer.

Code for producer process

```

while
while(true)
{
    /* produce an item in next-produced */
    while(counter == BUFSIZE);
    /* do nothing */
    buffer[in] = next-produced;
    in = (in + 1) % BUFSIZE;
    counter++;
}

```

Code for Consumer process

```

while(true)
{
    while(counter == 0);
    /* do nothing */
    next-consumed = buffer[out];
    out = (out + 1) % BUFSIZE;
    counter--;
    /* consume the item in next-consumed */
}

```

"Counter++" may be implemented in machine language

registers = Counter

registers = registers + 1

Counter = registers.

while registers is one of the local CPU registers.
Similarly, the statement "Counter--" is implemented as
follows

registers = Counter

registers = registers - 1

Counter = registers

Again registers is one of the local CPU registers.

The concurrent execution of "Counter++" and
"Counter--" is equivalent to a sequential execution in
which the lower-level statements presented
previously are interleaved in some arbitrary order

T ₀ :	producer execute	registers = counter	[registers = 5]
T ₁ :	producer execute	registers = registers + 1	[registers = 6]
T ₂ :	consumer execute	register ₂ = Counter	[register ₂ = 5]
T ₃ :	consumer execute	register ₂ = register ₂ - 1	[register = 4]
T ₄ :	producer execute	Counter = register ₂	[Counter = 6]
T ₅ :	consumer execute	Counter = register ₂	[Counter = 4]

Several processes access and manipulate the same data
concurrently and the outcome of the execution depends
on the particular order in which the access takes
place, is called a race condition.

do

[entry section]

Critical section

[exit section]

remainder section

z @while (true);

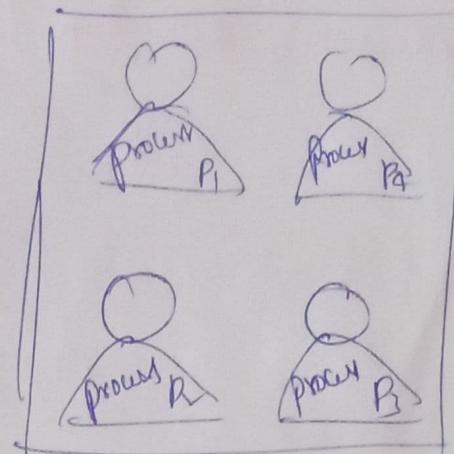
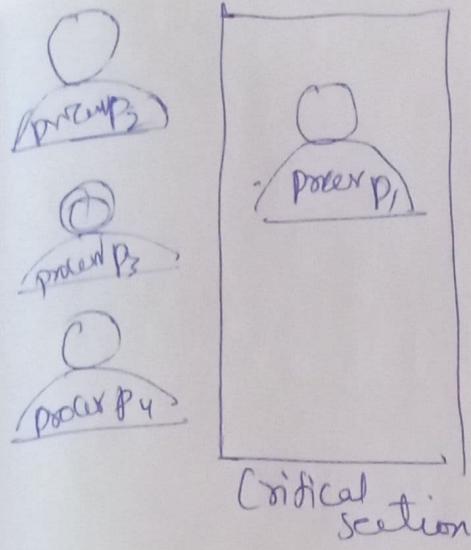
The Critical Section Problem

- The critical-section problem is to design a protocol that the processes can use to cooperate.
 - Each process must request permission to enter its critical section.
- ~~The section of code~~
- Critical section is the part of a program which tries to access shared resources.
 - The section of code implementing this request is the entry section.
 - The critical section may be followed by an exit section.
 - The remaining code is the remainder section.

A solution to the critical-section problem must satisfy the following 3 requirements:

① Mutual Exclusion

One solution must provide mutual exclusion. By mutual exclusion, we mean that if one process is executing inside critical section then the other process must not enter in the critical section.



progress! ✓ means that if one process doesn't need to execute into critical section then it should not stop other processes to get into the critical section

Bounded waiting : we should be able to predict the waiting time for every process to get into the critical section. the process must not be endlessly waiting for getting into the critical section.

Two general approaches are used to handle critical sections in operating systems?

preemptive kernel

nonpreemptive kernel

- ↗ A preemptive kernel allows a process to be preempted while it is running in kernel mode.
- ↗ A non preemptive kernel does not allow a process running in kernel mode to be preempted;

Synchronization Hardware

In synchronization hardware, we explore several more solutions to the critical section problem using techniques ranging from hardware to software based APIs available to application programmers.

These solutions are based on the premise of locking; however the design of such locks can be quite sophisticated.

→ These hardware features can make any programming task easier and improve system efficiency.

→ System efficiency decreases when this causes delays entry into each critical section &

→ Some modern computer systems therefore provide special hardware instructions that allow us that we can test and modify the content of a word or to swap the contents of two words atomically - i.e. as one interruptible unit.

The Test and Set() instruction can be defined as shown

boolean test_and_set (boolean *target)

{

 boolean sv = *target;
 *target = true;
 return sv;

}

Definition of test-and-set instruction.

The essential characteristic is that this instruction is executed atomically. So, if two TestAndSet instructions are executed simultaneously (each on different CPU), they will be executed sequentially in some arbitrary order.

We can implement mutual exclusion by declaring a boolean variable lock, initialized to false.

If the machine supports the TestAndSet() instruction. The structure of process P1 is shown below.

```
do {  
    while (test_and_set(&lock));  
    /* do nothing */  
    /* critical section */  
    lock = false;  
    /* remainder section */}
```

}

(b) Mutual-exclusion with while(true);
 definition of test-and-set() instruction

The swap() instruction, in contrast to get() and set() instruction, operates on the contents of two words; it is executed atomically.

* Mutual-exclusion can be provided as follows if the machine supports swap() instruction

→ Additionally, each process has a local boolean variable key.

```
int compare-and-swap(int *value, int expected, int newval)
    int temp = *value
    if (*value == expected)
        *value = newval;
    return temp;
```

↳ definition of compare-and-swap()

```
do
{
    while (compare-and-swap(&lock, 0, 1) != 0);
    /* do nothing */
    /* critical section */
    lock = 0;
    /* remainder section */
}
```

while (true);

mutual-exclusion implementation using
Compare-and-swap()

Since these algorithms satisfy the mutual-exclusion requirement, they do not satisfy the bounded-waiting requirement.

In below code, we present another algorithm using TestAndSet() instruction that satisfies all the critical-section requirements

do {

 waiting[i] = true;

 key = true;

 while (waiting[i] && key)

 key = test and set (&lock);

 waiting[i] = false;

 /* critical section */

 j = (i + 1) % n;

 while ((j != i) && !waiting[j])

 j = (j + 1) % n;

 if (j == i)

 lock = false;

 else

 waiting[i] = false;

 /* remainder section */

} while (true);

Bounded waiting mutual exclusion with test-and-set

Semaphore

def:

A semaphore is a shared integer variable after initialization we can only perform two operations on semaphore.

These are wait() and signal()

~~wait - P~~

The wait() operation was originally termed p (from the Dutch proverb, "to test"); signal() was originally called call v (from Verhogen, "to increment")

The definition of wait() is as follows

wait(s){

while ($s \leq 0$)

;|| no-op

$s--;$

spinlock

3.

The definition of signal() is as follows

signal(s){

$s++;$

3.

Two types of Semaphores

1. Binary Semaphore $s \geq 0/1$ - used for mutual exclusive locks
2. Counting Semaphore $s = 0, 1, 2, \dots$ to control access to a resource.
3. Semaphore mutex = 1;

do

{

 | wait(mutex);
 | Critical sections
 | Signal(mutex);

p_0, p_1, \dots, p_{n-1}

mutex

$\boxed{1/p_0}$

Semaphore printer=5;

$\boxed{\cancel{f} \cancel{f} \cancel{f} 3}$

 | remainder section
 | while(true);

 | wait(printer);
 | printeraccess p_0, p_2

 | signal(printer);

Semaphore $s \geq 0$

Synchronization

P1

:
:
:
S1
signal(synch),
:
:

P2

:
:
wait(synch),
S2
:
:

Explanation in next page

Usage

Binary Semaphores behave similarly to mutex locks.

→ It can be used instead for providing mutual exclusion.

Counting Semaphore can be used to control access to a given resource consisting of a finite number of instances.

Each process that wishes to use a resource performs a wait() operation on the semaphore (thereby decrementing the count) when a process releases a resource it performs signal() operation (incrementing the count).

We can also use semaphores to solve various synchronization problems.

For example) Consider two concurrently running processes P1 and P2 with a statement S₂

and P2 with a statement S₁.

Suppose we require that S₂ be executed only after S₁ has completed. We can implement this scheme readily by letting P2 and P1 share a common semaphore synch,

initialized to 0. But P2

P2 will execute S₂ only after P1 has invoked signal(synch), which after statement S₁ has been executed.

Semaphore producer-consumer problem

- The producer-consumer problem is a synchronization problem.
 - There's a fixed size buffer and the producer puts items and enters them into the buffer.
 - The consumer removes the items from the buffer and consumes them.
 - The producer consumer problem can be resolved using semaphores.
 - A producer should not produce items into the buffer when the consumer is consuming an item from the buffer and vice versa. So the buffer should only be accessed by the producer or consumer at a time.
- The producer and consumer processes share the following data structures:

```
int n;  
semaphore mutex = 1;  
semaphore empty = n;  
semaphore full = 0;
```

The mutex semaphore provides mutual exclusion for access to the buffer pool and is initialized to 1.

The empty and full semaphore count the number of empty and full buffer.

The semaphore empty is initialized to n. The semaphore full is initialized to 0.

The Code for the producer process.

```
do {  
    /* produce an item in next-produced */  
    ...  
    wait(empty);  
    wait(mutex);  
    /* add next-produced to the buffer */  
    ...  
    signal(mutex);  
    signal(full);  
} while(true);
```

The Code for the consumer process.

```
do {  
    wait(full);  
    wait(mutex);  
    /* remove an item from buffer to next-consumed */  
    ...  
    signal(mutex);  
    signal(empty);  
    /* consume the item in next-consumed */  
    ...  
} while(true);
```

We can interpret the code as the producer producing full buffers for the consumer or as the consumer producing empty buffers for the producer.

Reader-writer problem

The reader-writer problem relates to an object such as a file that is shared by multiple programs.

Some of the processes are readers, they want to read the data from the object.

Some of the processes are writers, they want to write into the object.

If two readers access the object at the same time there is no problem.

However, if two writers or a reader and writer access the object at same time there may be problems.

To solve this situation, a writer should get exclusive access to an object. This can be implemented using semaphores.

In the solution to the first reader writer problem, the reader process share the following data structures:

```
Semaphore rw-mutex=1; } initialized to 1  
Semaphore mutex=1;  
int read-count=0; → initialized to 0.
```

~~mutex~~ and

Writer process

The code that defines the writer process is given below:

```
do {
```

```
    wait (rw-mutex);
```

```
    ...
```

/* writing is performed */

```
    ...
```

```
    Signal (rw-mutex);
```

```
    while (true);
```

Dining-philosop

If a writer wants to access the object, wait operation is performed. After that no other writer can access the object. When a writer is done writing into the object, Signal operation is performed.

Reader process

The code that defines the reader process is given below:

```
do
    wait(mutex);
    read_count++;
    if(read_count == 1)
        wait(rw_mutex);
    signal(mutex);
    ...
    /* reading is performed */
    ...
    wait(mutex);
    read_count--;
    if(read_count == 0)
        signal(rw_mutex);
    signal(mutex);
}
while(true);
```

* The mutex semaphore ensures mutual exclusion and rw-mutex handles the writing mechanism.

* The variable ~~read~~ denotes read_count denotes the no. of readers accessing the object.

The var to the mutex is

As soon as read_count becomes 1, wait operation is used on rw-mutex. This means that a writer cannot access the object anymore.

After the read operation is done, read_count is decremented.

When read_count becomes 0, signal operation is used on rw-mutex so a writer can access the object now.

Dining philosophers problem

Consider 5 philosophers who spend their lives thinking.
There are 5 philosophers sharing a circular table
and they eat and think alternatively.

There is a bowl of rice for each of the philosophers
and 5 chopsticks.

A philosopher needs both their right and left
chopstick to eat. A hungry philosopher may only
eat if there are both chopsticks available.
Otherwise a philosopher puts down their chopstick
and begins thinking again.

The dining-philosophers problem is considered
a classic synchronization problem neither
because of its practical importance nor because
computer scientists dislike philosophers.

→ It demonstrates a large class of concurrent
control problems.

→ It is a simple representation of the need
to allocate several resources among several
processes in a deadlock-free and starvation-
free manner.

One simple solution is to represent each
chopstick with a semaphore.

A philosopher tries to grab a chopstick by
executing wait() operation on the semaphore.
She releases her chopstick by executing the signal
operation.

Thus, the shared data is
semaphore chopstick[5];

```
do {  
    wait(chopstick[i]);  
    wait(chopstick[(i+1)%5]);  
    ...  
    /* eat for a while */  
    signal(chopstick[i]);  
    signal(chopstick[(i+1)%5]);  
    ...  
    /* think for a while */  
    ...  
} while(true);
```

The above solution makes sure that no two neighboring philosophers can eat at the same time. But this solution lead to a deadlock.

This may happen if all the philosophers pick their left chopstick simultaneously. Then none of them can eat and deadlock occurs.

Some of the ways to avoid deadlock are as follows -

- There should be at most four philosopher on the table.

Monitors

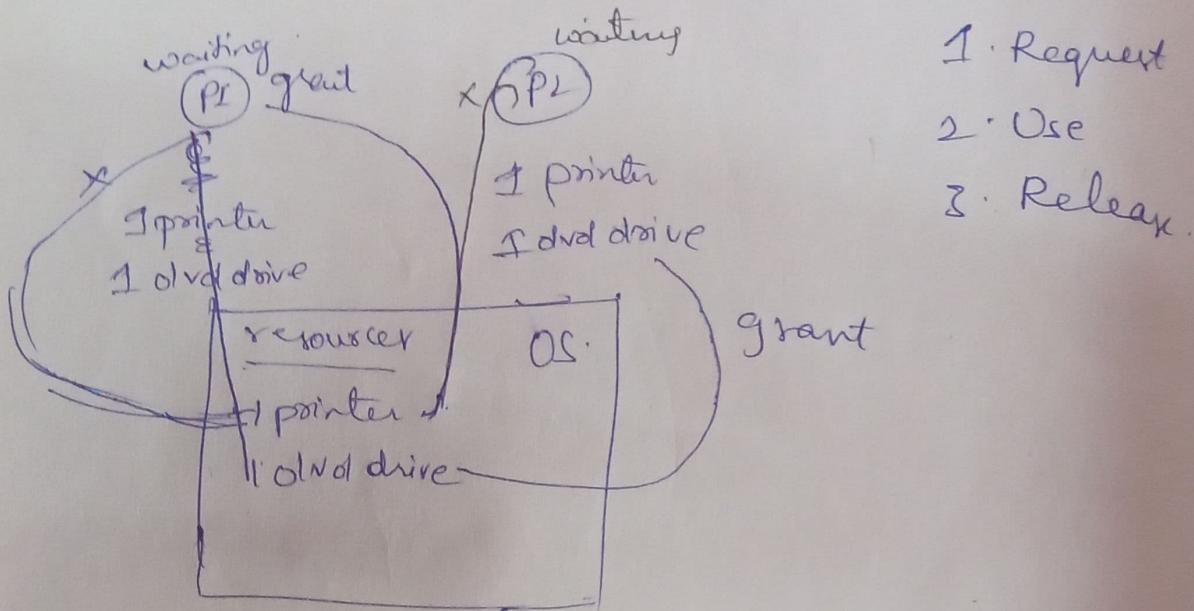
- An even philosopher should pick the right chopstick and then the left chopstick, while an odd philosopher should pick the left chopstick and then right chopstick.
- A philosopher should only be allowed to pick their chopstick if both are available at the same time.

Monitors

Dead locks.

Resource request

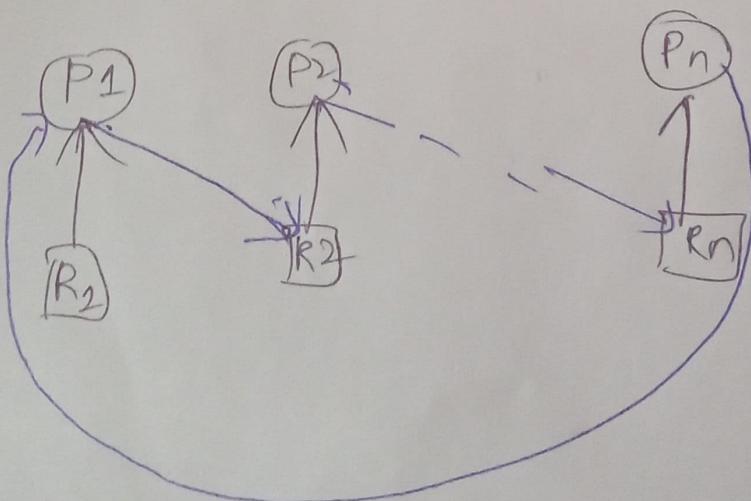
1. Available - it is allocated to process.
2. Not Available - it enters into blocked.



1. Request
2. Use
3. Release

Necessary Condition for dead lock .

1. Mutual Exclusion.
2. Hold and wait
3. No preemption
4. Circular wait



System Resource Allocation graph

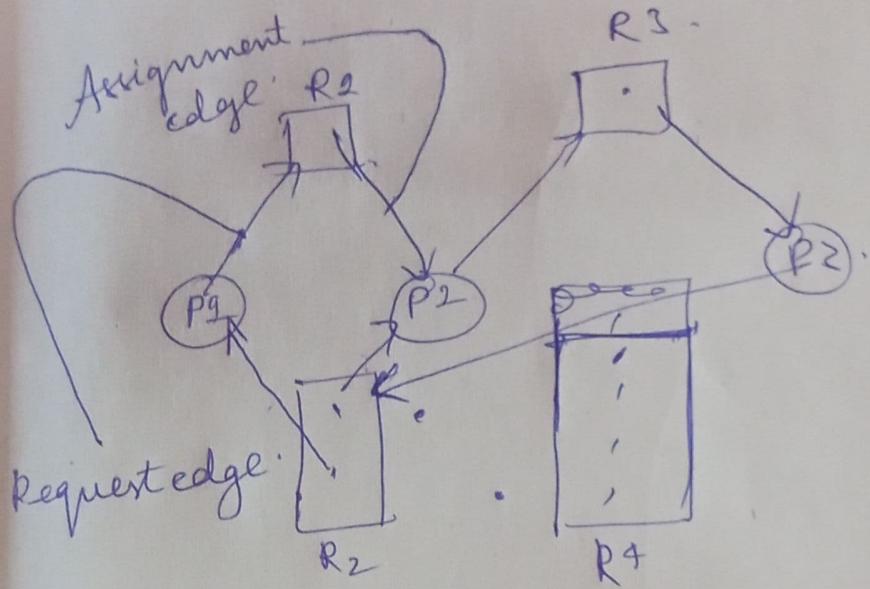
$$G = (V, E)$$

V is non empty set of vertices.

E is pair of vertexies called edge.

v - process - circle

resources - rectangle



If Resource Allocation graph containing cycle (closed path)
there is a possibility of deadlock.

$$P_2 \rightarrow R_3 \rightarrow P_3 \rightarrow R_2 \rightarrow P_2$$

$$P_2 \rightarrow R_3 \rightarrow P_3 \rightarrow R_2 \rightarrow P_1 \rightarrow R_1 \rightarrow P_2$$

Each resource have one instance, and having cycle. In that case, compulsory deadlock is present (or) exists

Resources multiple instance - graph contains cycle \rightarrow deadlock may exists.

~~cycle is necessary~~